

Fire Detection With Erroneous Input Correction

ECGR 4105/5105 - Introduction to Machine Learning

1st Jaskin Kabir
ECGR Student
UNC Charlotte
Charlotte, NC
jkabir@charlotte.edu

2nd Axel Leon Vasquez
ECGR Student
UNC Charlotte
Charlotte, NC
aleonvas@charlotte.edu

3rd Matthew Anderson
ECGR Student
UNC Charlotte
Charlotte, NC
mande137@charlotte.edu

Abstract—The accuracy of typical fire alarm systems is often limited due to their reliance on single sensors with constrained detection capabilities. Moreover, a sensor failure renders the entire system ineffective, posing a critical risk. While AI-based fire detectors have massively outperformed traditional systems, little work has addressed handling sensor failures. To address this challenge, we propose AlarmNet, a robust fire detection system that incorporates erroneous sensor data during training. We investigate two implementations of AlarmNet: a centralized model trained on a single server and a simulation of a federated model where training is distributed across multiple clients and aggregated into a global model. In our experiments, sensor errors were introduced by randomly replacing 24.2% of the dataset’s measurements with missing or faulty values. We applied a pre-processing step that replaced these values with the median of the corresponding training features. Our 3-layer artificial neural network The model achieved an accuracy of 98.38% and successfully detected 99.73% of fires, despite the high rate of sensor errors. This resulted in only a 0.25% reduction in detection performance compared to a model trained on complete data, and a significant 41.51% improvement over a simulated commercial-grade fire alarm system.” The federated approach delivered comparable results while utilizing a less computationally intensive model. In future work, we aim to extend this approach to develop a real-world fire detection system capable of maintaining robustness against sensor failures.

Index Terms—machine learning (ML), algorithm, model, Artificial Neural Network, Imputation, Sensor fusion, Fire Alarms, Fire Detection, Error Handling, Federated Learning

I. INTRODUCTION AND MOTIVATION

Fire alarm systems play a critical role in safeguarding lives and property by providing early warnings of fire hazards. Traditional systems, however, do not perform as well as their critical role would suggest. For example, a 2013 study found the Honeywell FS90, a widely used commercial fire detection system, only achieves an accuracy of 87.5% [1]. This contributes to deaths, injuries, and damage due to fires in two ways. (1) When a fire occurs and the fire detection system fails to recognize it, the occupants of the building can only react to the fire once they are close enough to observe it, often once it is too late to evacuate. (2) A 1995 report stated that fewer than 25% of people interpreted the sound of the fire alarm as a potential indication of a real emergency [2]. Due to the high frequency of false or ‘nuisance’ alarms, occupants may not

evacuate or contact emergency services, even if the alarm is triggered by a real fire.

As these systems typically rely on just one or two sensors and simple logic to detect fires, they are especially susceptible to faulty sensors rendering them completely ineffective. While artificial intelligence (AI)-based fire detectors have demonstrated significant improvements in accuracy over traditional systems, they generally rely on error-free sensor data [3] [4] [5]. This limits their practical applicability, especially in real-world scenarios wherein sensor errors and failures are common.

To address these challenges, we propose AlarmNet, a robust, AI-based fire detection system specifically engineered to remain accurate in the presence of sensor errors. AlarmNet is trained on a dataset from which randomly selected measurements were removed to simulate missing or faulty sensor readings. This approach enables the model to learn to detect fires accurately based on unreliable sensor data.

The system is evaluated in two configurations: a centralized model trained on a single server and a simulation of a federated model, where training is distributed across multiple clients and aggregated into a global model. This dual approach allows us to investigate two key questions: (1) whether it is possible to train a model that can handle sensor errors, and (2) whether a federated model can achieve comparable performance while utilizing fewer computational resources. If the federated model meets this goal, it would demonstrate AlarmNet’s viability as a real-world fire detection system. This paper first details the analysis and preparation of the chosen dataset, followed by a discussion of the design, implementation, and evaluation of both versions of AlarmNet.

II. APPROACH

III. DATA PREPARATION

A. Dataset

The dataset used in this project is the Smoke Detection Dataset provided by Stefan Blattmann [6]. The dataset consists of approximately 62,629 sensor readings encompassing both normal and fire-related scenarios, providing diverse training and evaluating fire detection models.

To generate this data, Blattman constructed an Arduino-based system to collect data from an array of six sensors at a 1Hz sampling rate. These sensors were:

- Bosch BMP390: Pressure Sensor
- Bosch BME688: Humidity, Temperature, and Gas Sensor (Volatile Organic Compounds/VOC)
- Bosch BMP388: Pressure Sensor
- Sensirion SPS30: Particulate Matter Sensor
- Sensirion SHT31: Humidity and Temperature Sensor
- Sensirion SPG30: Gas sensor (VOC)

Note that several sensors measure the same environmental factors, such as humidity and pressure. This redundancy is intentional and serves to enhance the robustness of the model. Before uploading this dataset to Kaggle, Blattman utilized sensor fusion to combine the data from these six sensors into a single dataset. The only sensor with no redundancy is the Sensirion SPS30, which measures particulate matter and is the most expensive sensor.

Blattman placed his sensor array into various fire scenarios and recorded the measurements. He provided the following short list of scenarios during which he collected data:

- Normal indoor
- Normal outdoor
- Indoor wood fire, firefighter training area
- Indoor gas fire, firefighter training area
- Outdoor wood, coal, and gas grill
- Outdoor high humidity
- etc.

[7]

B. Train/Test Split

Before the dataset was passed to the models, it was split using an 80/20 train/test split. This means that 80% of the data were randomly sampled and used to train the model, while the remaining 20% were used to evaluate the model's performance. This split ensures that the model is not overfitting to the training data, as it is evaluated on data that it has not seen before.

C. Feature Scaling

After the train/test split, the features were scaled using standard scaling. This method subtracts the mean of the feature and divides by the standard deviation, ensuring that all features have a mean of 0 and a standard deviation of 1. This scaling is important because it ensures that the model does not learn to prioritize features with larger values, which could lead to poor generalization. The mean and standard deviation are calculated based on the training data and then applied to the test data to ensure that the model is evaluated on the same scale that it was trained on.

D. Feature Elimination

The dataset provided by Blattman originally included 15 features, which are listed in Table I. Of these 15, the two sample count features and the timestamp were removed as they do not provide any useful information for the model. Of

the remaining 12 features, a crucial task was to determine which features were the most relevant to the target variable. Eliminating irrelevant variables reduces the complexity and computational cost of the model, as its input space shrinks in dimensionality. This reduction can also help prevent overfitting, as the model is less likely to learn noise in the data.

Feature	Description
Timestamp	UTC Timestamp of the sample
Number	Unique identifier for each sample
Unnamed: 0	Unintended duplicate of Number
Temperature	Air temperature in degrees Celsius
Humidity	Air humidity in percentage
Pressure	Air pressure in hectoPascals
TVOC	Total Volatile Organic Compounds in parts per billion
eCO2	Carbon Dioxide in parts per million
Raw H2	Raw, uncompensated molecular hydrogen
Raw Ethanol	Raw, uncompensated ethanol
PM1.0	Percentage of particles in the air less than $1.0\mu m$ in diameter
PM2.5	Percentage of particles in the air less than $2.5\mu m$ in diameter
NC0.5	Number concentration of particles in the air greater than $0.5\mu m$ in diameter
NC1.0	Number concentration of particles in the air greater than $1.0\mu m$ in diameter
NC2.5	Number concentration of particles in the air greater than $2.5\mu m$ in diameter

TABLE I
SMOKE DETECTION DATASET FEATURES

To achieve this, we calculated the absolute correlation of each feature with the target variable, 'Fire Alarm'. The features were then ranked based on this correlation, and the top four features were selected for the model. This number was chosen by incrementally removing the least correlated features from the dataset and training a model on the these data until a significant drop in accuracy was measured. These top four features were:

- 1) Humidity
- 2) Raw Ethanol
- 3) Pressure
- 4) TVOC

While a sufficiently accurate model could be trained on these data, this list of features does not include any of the particulate matter measurements. If this system were to ignore one of its sensors, it would be significantly less resilient against sensor errors. To address this, we added back the NC0.5 feature to the data, as it is the most correlated with the target variable among the particulate matter measurements.

E. Error Insertion

To simulate sensor errors, we developed a method to randomly remove certain measurements from the dataset accord-

ing to some assumed rate of error. This rate of error represents the chance that a given sensor reading will provide incorrect or missing data during any given sample. We chose to assume this error rate to be 40%. This is an unrealistically high number, but it is low enough to ensure that the model is not completely useless while still providing a significant challenge.

Most of the features—with the exception of the particulate matter measurements—were collected using two sensors. To include this redundancy in the simulation, we assumed that both sensors that generate any feature would have to fail simultaneously in order for that feature to be missing for a given a sample. In other words, while 40% of the particulate matter measurements were removed, only 20%¹ of the other measurements were removed. After applying this error simulation to the dataset, approximately 24.2% of the data was removed.

F. Imputation

Before this data could be used to train a model, the missing values had to be replaced. We chose to use median imputation, which replaces missing values with the median of the corresponding feature calculated from the training dataset. This method was chosen because it is robust to outliers and skewed distributions, which are common in sensor data.

IV. MODEL TRAINING AND EVALUATION

A. Evaluation Methods

1) Multiple Models

To ensure that the techniques used in this project are effective, we trained and evaluated multiple models. We trained a model based on the full 12-features of the dataset, a model based on the top 4 features, a model that included the NC0.5 feature for redundancy. We then compared the performance of these three models to determine if the feature elimination was necessary, and how adding back the NC0.5 feature affected this performance.

We also trained a model on the imputed error data, and a simulation of a federated learning system. The performance of the federated model was then compared to the centralized model.

2) Metrics

The performance of each model was evaluated using the following metrics:

- **Recall:** Measures the proportion of actual fire cases detected by the model to minimize false negatives. This is the most important metric, as it is crucial that the model detects as many fires as possible.

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (1)$$

- **Precision:** The proportion of correctly identified fire alarms

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \quad (2)$$

¹The joint probability formula $P(A \cap B = P(A) \times P(B))$ implies that this probability should actually be 16%. We chose to use 20% instead for the added challenge.

- **F1-Score and Accuracy:** Implements the harmonic mean of precision and recall with overall correctness of the model.

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} \quad (3)$$

- **Confusion Matrix:** A visualization tool that lists the true positive/negative and false positive/negatives

3) Honeywell FS90 Simulation

The 2013 study mentioned in the introduction only listed the Honeywell FS90's accuracy with no mention of precision, recall, or the raw data necessary to calculate these values [1]. To provide a more accurate comparison, we developed a simple simulation of the FS90:

First, we created a copy of the target column from the training data. Then we randomly selected 12.5% of the data-points in the data to flip, creating the FS90's 87.5% accurate 'prediction'. To simulate sensor error, we assumed that the FS90 had a 20% chance of experiencing a sensor error during any given sample. This error rate was chosen to match the rate of error in the AlarmNet model. We then assumed that during a sensor error, the FS90 would indicate that it did not detect a fire. To apply this to the simulation, we randomly selected 20% of the prediction data. If the FS90 predicted a detected a fire during any of these samples, we flipped the prediction to indicate that it did not detect a fire. This simulation allowed us to calculate the recall and precision of the FS90 with and without sensor error, which we then compared to the AlarmNet models.

4) SVM Benchmark

We trained a Support Vector Machine(SVM) on the same dataset as the model trained on the imputed error data. The SVM model was chosen as it was the best performing of the classical machine learning techniques on this dataset. The performance of this model was compared to the fully connected models to ensure that neural networks are the necessary for this application.

B. Model Architectures

For the model trained on the 12-input model, the fully connected neural network included an input layer of 12 neurons, two hidden layers of 64 and 32 neurons respectively, and a layer consisting of one output neuron. The 4 and 5-input models used input layers of 4 and 5 neurons respectively, two hidden layers both consisting of 64 neurons, and one output neuron. Between each layer, a ReLU activation function was used. This introduced nonlinearity into the model, allowing it to learn complex patterns in the data and ensuring its loss function was differentiable. After the output layer, a sigmoid activation function was applied to ensure that the output was between 0 and 1.

To allow the centralized model trained on the data with errors inserted into it, the architecture was increased to a model with three hidden layers each consisting of 256 neurons. This allowed the model to learn more complex patterns in the data and ensured that it could handle the noisy and incomplete data.

The model used in the federated learning simulation used three hidden layers consisting of 32, 16, and 8 neurons. This model includes fewer neurons than any other model to ensure that it is less computationally intensive. This model was trained on three clients, each with a copy of the global model. The clients trained their local models on their own data, and then sent the weights of their models to the global model. The global model then averaged these weights to create a new global model.

C. Training

1) Network Classes

To train these models, we used the PyTorch library to develop the `AlarmNet` class. This class has a configurable number and size of hidden layers so that a new class doesn't have to be written for each network configuration. The class also includes methods to train the model, evaluate its performance, and generate visualizations of the model's performance. The class was then used to train the models on the training data and evaluate their performance on the test data.

Additionally, a `FederatedLearning` class was developed to simulate the federated learning system. Upon instantiation, this class instantiates a global instance of the `AlarmNet` class. During training, it splits the dataset into n parts and then instantiates n copies of the global `AlarmNet` class to represent n different clients. Each of these classes is then trained on a different part of the dataset. After training, the weights of each model are averaged to update the global model. This model is then evaluated on the test data to determine its performance. n is a configurable parameter, and so are the number of training epochs each client trains for and how many times this process is repeated.²

2) Training Parameters

The models were trained using the Adam optimizer. The loss function used was binary cross-entropy, which is commonly used for binary classification problems. The models trained on the data without inserted errors used a learning rate of 10^{-3} and were trained for 3000 epochs.

To accommodate the error model's larger architecture, the learning rate was reduced to 10^{-5} and the model was trained for 16000 epochs.

The federated model was trained with a learning rate of 10^{-3} for the global model and the clients. We chose to simulate three clients, each performing 600 training epochs per round. The process of training the client models for 600 epochs and averaging their weights was repeated 3 times.

V. RESULTS AND ANALYSIS

The model trained on the full 12-features achieved a near 100% accuracy, which can be seen in the confusion matrix depicted in Figure 4. While this model achieved high precision and recall, its computational complexity makes it less inefficient, and difficult to scale for handling the errored data.

The model trained on the top four most correlated features did not perform as well, however. As shown in Figure 6,

eliminating all but the top four features lead to a drop in accuracy of 0.07% and a drop in precision of 0.1%. While this drop in performance is minimal, it is still significant. This suggests that the particulate matter measurements are important for the model's performance.

This is further supported by the confusion matrix displayed in Figure 11 and the metrics shown by Figure 9. By adding back the NC0.5 feature, the model was able to achieve almost the exact same performance as the 12-input model.

After applying the error insertion and imputation, the model still performed impressively well. As shown in Figure 12, the model trained on the errored data achieved an accuracy of 98.36% and a recall of 99.76%. This marks a drop in accuracy of 1.66%, but a drop in recall of only 0.21%. As recall is the most important metric for a fire detection system, this performance is much higher than expected. The confusion matrix for this model can be seen in Figure 14.

As a benchmark, the SVM model was trained on the same training dataset and evaluated on the same validation data as the model trained on the errored dataset. From the metrics shown in Figure 1, This model was 6.14% less accurate than the centralized model, and scored 0.26% lower in recall. This demonstrates that the neural network is better suited for this application than the SVM.

SVMs are known for their simplicity and effectiveness within smaller datasets due to their ability to model linear and non-linear relationships using kernel functions. However, in this dataset, the project had a large size of 62,000 samples which limited the SVM as its memory requirement scales quadratically with the number of samples. On the other hand the neural networks demonstrated superior performance and adaptability. They efficiently scaled with large datasets by leveraging GPU acceleration with mixed precision and gradient scaling.

As shown in Figure 15 The 12-input federated model performed worse than the 12-input centralized model, likely due to its reduced complexity. However, the 5-input federated model performed almost identically to the 5-input centralized model, which can be seen in Figure 16.

Additionally, Figure 17 shows that the federated model trained on the errored dataset still performs exceptionally well in terms of precision and recall.

Finally, Figure 19 and Figure 18 show a comparison between the federated and centralized models, as well as the Honeywell FS90 in terms of precision and recall respectively. The federated model performed almost identically to the centralized model and both massively outperformed the FS90. This shows that the federated model was able to achieve incredible accuracy and recall while being the least computationally intensive model and proves that the AlarmNet system is a viable solution for fire detection.

VI. LESSONS LEARNED

This machine learning model successfully demonstrated the development an efficient fire detection system using sensor data. By leveraging a custom neural network architecture, the

²The code can be found at https://github.com/jaskinkabir/Intro_ML_Project

system achieved high precision and recall even with real-world scenarios such as noisy or missing sensor data. Feature selections and target preparation, including median imputation and error handling, ensure the model's exploration. The reduced feature model maintained comparable performance to the 12-input model significantly reducing computational requirements and making it good for constrained environments. Comparative analysis highlights the proposed neural network over the traditional such as SVM. The neural network can handle large datasets, capture complex features tolerate sensor errors underscores the effects in IoT-base fire detection systems. Optimizing time and compact model size ensure viability for read-time devices. The project object results emphasize the importance of designing models that are balanced, efficient, accurate, and robust for safety and critical protocol applications. Future work can be explored further optimizing the learning, advancing imputation, and improving the federated model weights between three clients to enhance system performance. The project demonstrates the potential for machine learning to revolutionize fire detection technology to challenge high-end companies such as Honeywell to ensure the power of ANN to contribute to improving safety and reliability in diverse environments.

VII. CONTRIBUTIONS

- **Matthew Anderson:**
 - Federated Learning Class
 - Edited Report and Presentation
- **Jaskin Kabir:**
 - Project Manager
 - Dataset Selection, Preparation, and Analysis
 - AlarmNet Class that contained the centralized model
 - Error insertion and imputation
 - Edited Report and Presentation
- **Axel Leon Vasquez:**
 - Wrote Report
 - Created Presentation
 - Data Visualization

REFERENCES

- [1] N. k. Fong, "Investigation of the performance and improvement of optical smoke detectors," *Procedia Engineering*, vol. 62, 12 2013.
- [2] G. Proulx, J. C. Latour, J. W. MacLaurin, J. Pineau, L. E. Hoffman, and C. Laroche, "Housing evacuation of mixed abilities occupants in highrise buildings," National Research Council of Canada. Institute for Research in Construction, Tech. Rep., Aug 1995. [Online]. Available: <https://doi.org/10.4224/20375582>
- [3] H. Prasad, A. Singh, J. Thakur, C. Choudhary, and N. Vyas, "Artificial intelligence based fire and smoke detection and security control system," in *2023 International Conference on Network, Multimedia and Information Technology (NMITCON)*, 2023, pp. 1–6.
- [4] D. K. Dewangan and G. P. Gupta, "Explainable ai and yolov8-based framework for indoor fire and smoke detection," in *2024 IEEE International Conference on Information Technology, Electronics and Intelligent Communication Systems (ICITEICS)*, 2024, pp. 1–6.
- [5] H. Zhong, Z. Wang, Z. Chen, W. Chen, and Y. Li, "A novel fire monitoring system for electric bicycle shed based on yolov8," in *2023 IEEE 16th International Symposium on Embedded Multicore/Many-core Systems-on-Chip (MCSoc)*, 2023, pp. 142–147.
- [6] S. Blattman, "Smoke detection dataset version 1," Aug 2022. [Online]. Available: <https://www.kaggle.com/datasets/deepcontractor/smoke-detection-dataset>
- [7] S. Blattmann, "Real-time smoke detection with ai-based sensor fusion," Aug 2022. [Online]. Available: <https://www.hackster.io/stefanblattmann/real-time-smoke-detection-with-ai-based-sensor-fusion-1086e6#toc-background-1>

Figures

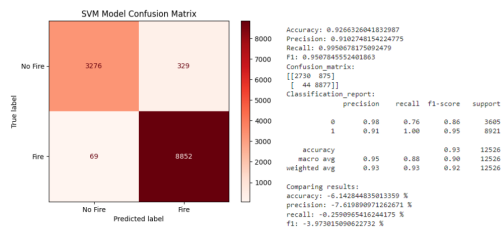


Fig. 1. SVM Metric

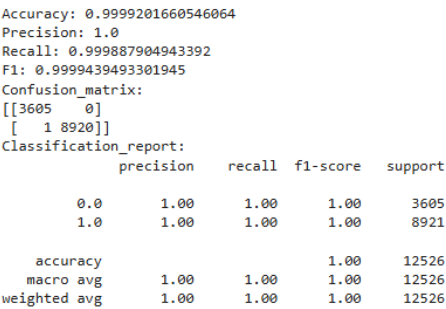


Fig. 2. 12 - Feature Metrics




Fig. 3. 12 - Feature Loss Curve




Fig. 4. 12 - Feature Confusion Matrix

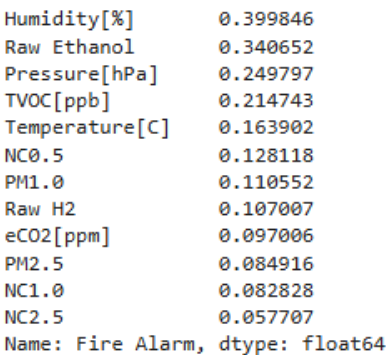


Fig. 5. Features

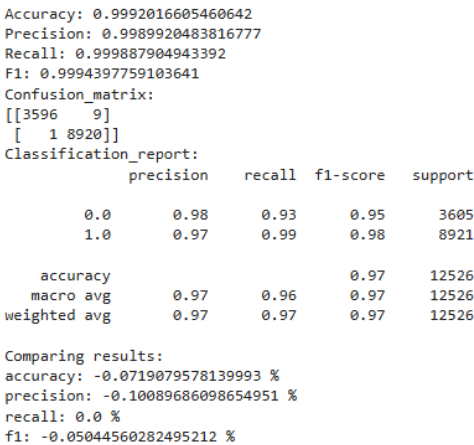


Fig. 6. 4 - Feature Metrics




Fig. 7. 4 - Feature Loss Curves




Fig. 10. 5 - Feature Loss Curves




Fig. 8. 4 - Feature Confusion Matrix




Fig. 11. 5 - Feature Confusion Matrix

```

Accuracy: 0.9998403321092129
Precision: 1.0
Recall: 0.9997758098867839
F1: 0.9998878923766816
Confusion_matrix:
[[3605  0]
 [ 2 8919]]
Classification_report:
      precision    recall  f1-score   support

     0.0         1.00      1.00      1.00       3605
     1.0         1.00      1.00      1.00       8921

   accuracy          1.00      1.00      1.00      12526
  macro avg          1.00      1.00      1.00      12526
 weighted avg          1.00      1.00      1.00      12526

Comparing results:
accuracy: 0.06387735547748555 %
precision: 0.10079516183223447 %
recall: -0.011212019284675844 %
f1: 0.044816670922210436 %

```

Fig. 9. 5 - Feature Metrics

```

Accuracy: 0.9835542072489223
Precision: 0.9796367638965328
Recall: 0.9976460038112319
F1: 0.9885593690991892
Confusion_matrix:
[[3420 185]
 [ 21 8900]]
Classification_report:
      precision    recall  f1-score   support

     0.0         0.98      0.97      0.98       3605
     1.0         0.99      0.99      0.99       8921

   accuracy          0.99      0.97      0.98      12526
  macro avg          0.99      0.97      0.98      12526
 weighted avg          0.98      0.98      0.98      12526

Comparing results:
accuracy: -1.6558441558441537 %
precision: -2.078651685393257 %
recall: -0.213483146067411 %
f1: -1.1459678659243208 %

```

Fig. 12. Imputation Metrics




Fig. 13. Imputation Loss Curves




Fig. 17. Federated Learning: Error handling




Fig. 14. Imputation Confusion Matrix

0.999775808867839
0.9976460038112319
(np.float64(1.0), np.float64(0.9607843137254902), np.float64(0.9411764795882353), np.float64(1.0))

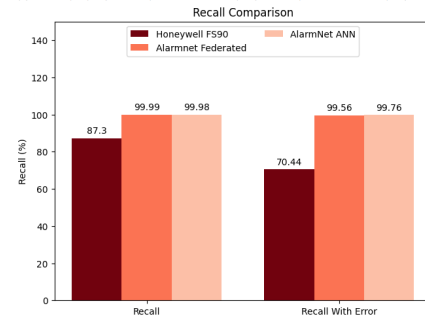


Fig. 18. Recall Comparison Honeywell FS90 Vs AlarmNet ANN Vs Federated

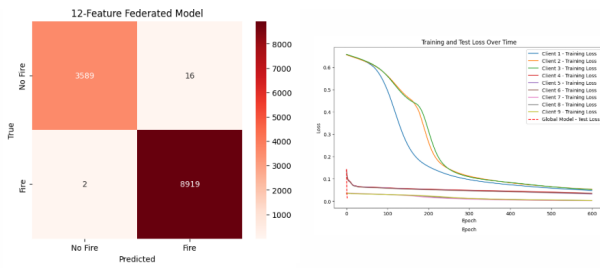


Fig. 15. Federated Learning: 12-Features




Fig. 16. Federated Learning: 5-Features




Fig. 19. Precision Comparison Honeywell FS90 Vs AlarmNet ANN Vs Federated